

Bases de Datos

Clase 10: Evaluación de consultas

Páginas, disco y buffer

Para trabajar con las tuplas de una relación, la base de datos carga la página desde el disco con dicha tupla

Para cargar estas páginas, la base de datos reserva un espacio en RAM llamado **Buffer**

Costo de un algoritmo

Cuántas veces tengo que leer una página desde el disco, o escribir una página al disco!

Las operaciones en buffer (RAM) son orden(es) de magnitud más rápidas que leer/escribir al disco – costo 0

Algoritmos en una BD

Implementan interfaz de un iterador lineal:

- open()
- next()
- close()

Algoritmos en una BD

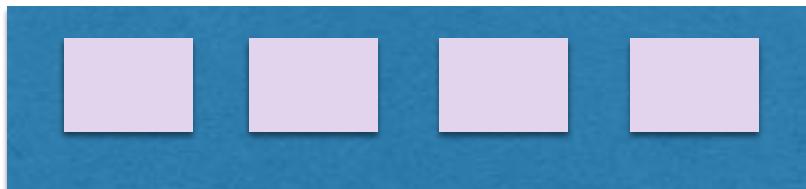
Para una relación R:

- R.open() – se posiciona **antes** de la primera tupla de R
- R.next() – devuelve la siguiente tupla o NULL
- R.close() – cierra el iterador

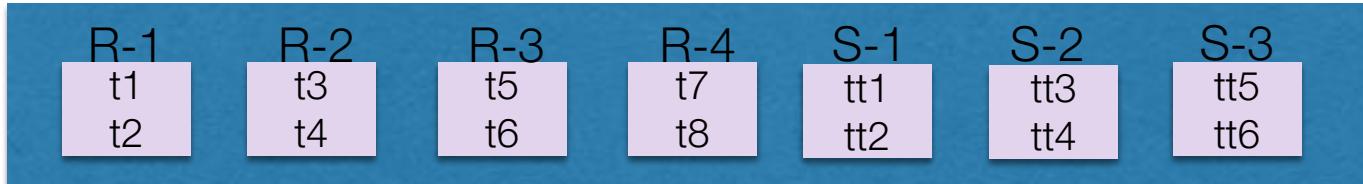
Páginas, disco y buffer

DB

Buffer



Disco

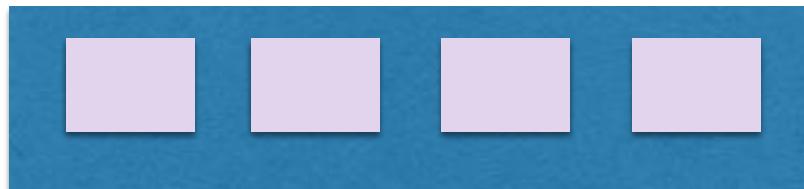


Páginas, disco y buffer

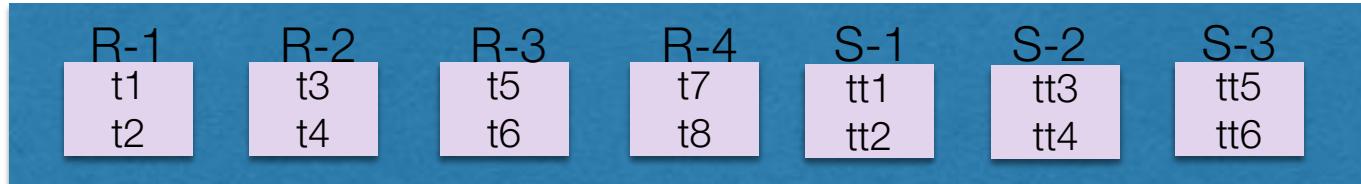
DB

R.open()

Buffer



Disco

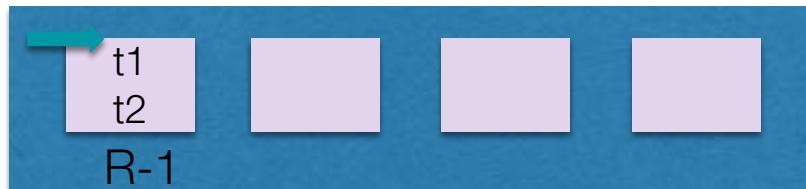


Páginas, disco y buffer

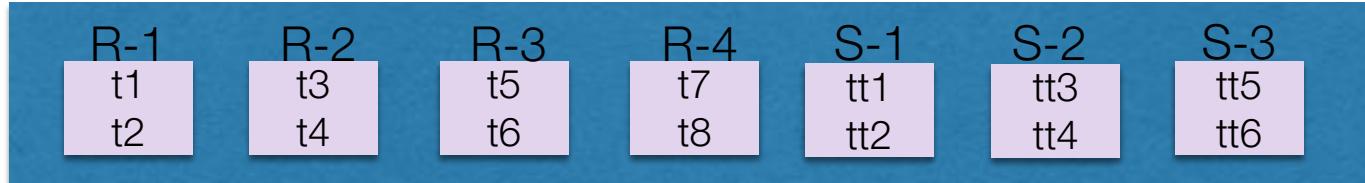
DB

R.open()

Buffer



Disco

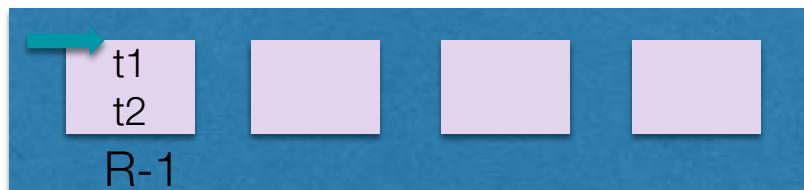


Páginas, disco y buffer

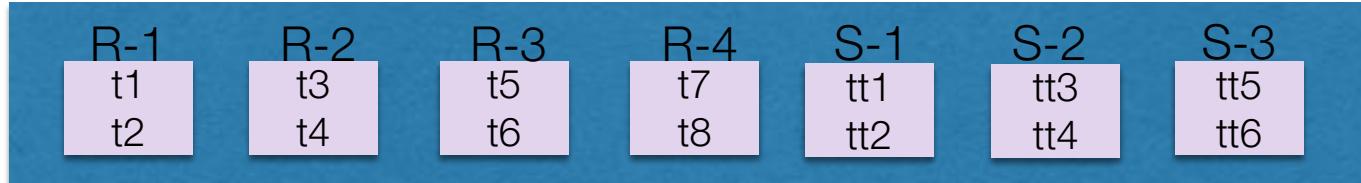
DB

R.next()

Buffer



Disco

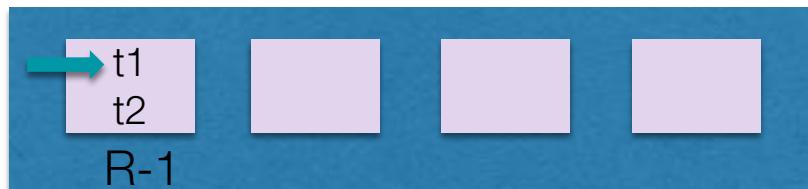


Páginas, disco y buffer

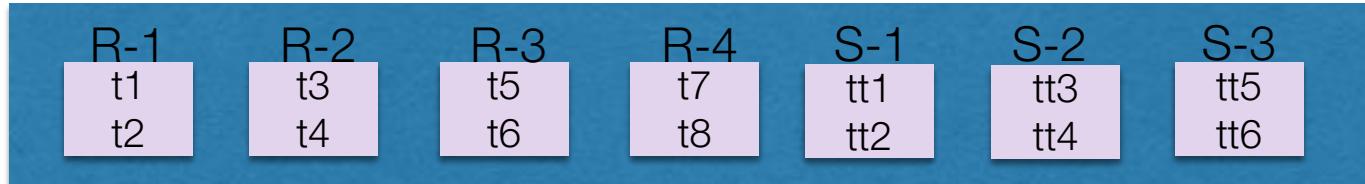
DB

R.next()

Buffer



Disco

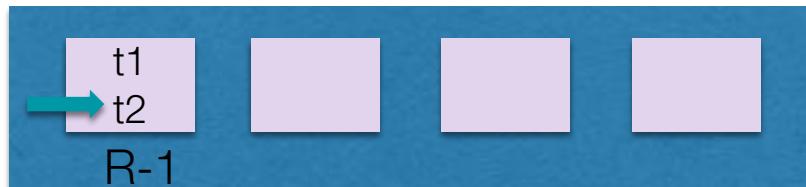


Páginas, disco y buffer

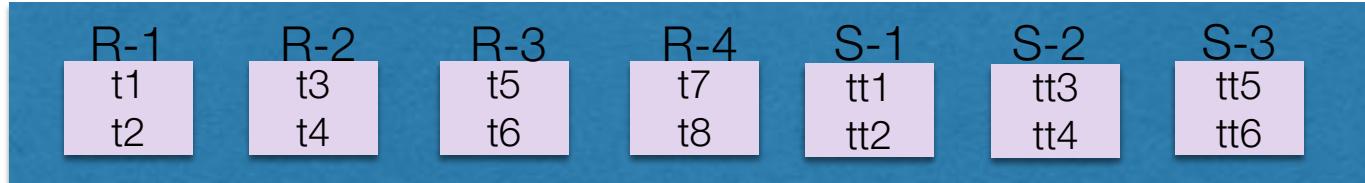
DB

R.next()

Buffer



Disco

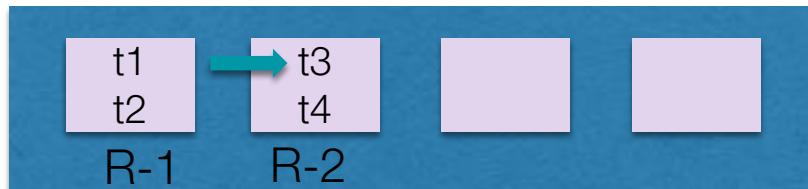


Páginas, disco y buffer

DB

R.next()

Buffer



Disco



Páginas, disco y buffer

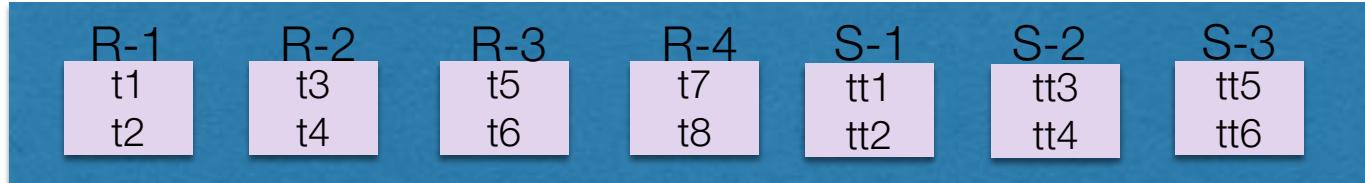
DB

R.close()

Buffer



Disco

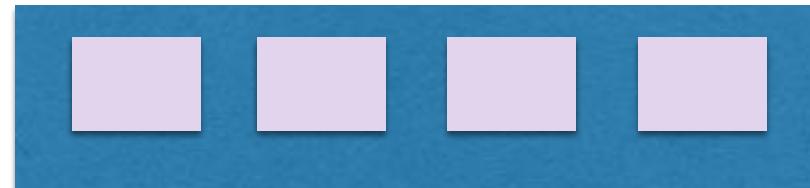


SELECT * FROM R

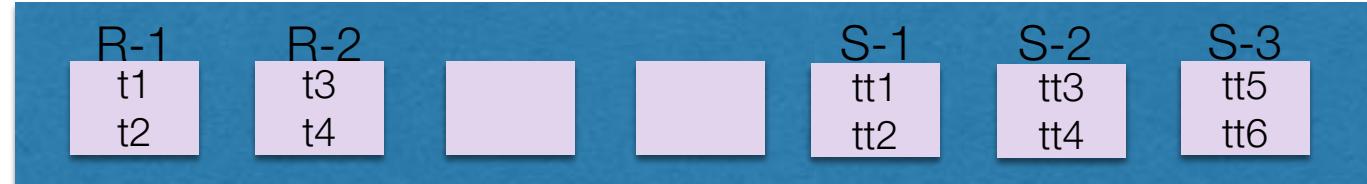
DB

Buffer

Disco



```
R.open()  
t:= R.next()  
while t != null do  
    output t  
    t:= R.next()  
  
R.close()
```

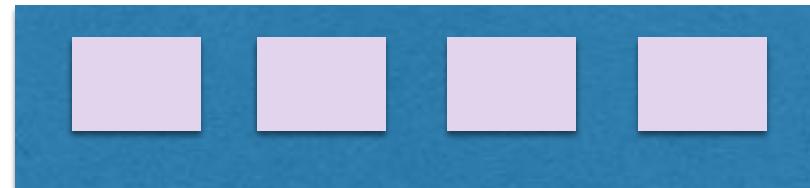


SELECT * FROM R

DB

Buffer

Disco



R.open()

```
t := R.next()  
while t != null do  
    output t  
    t := R.next()
```

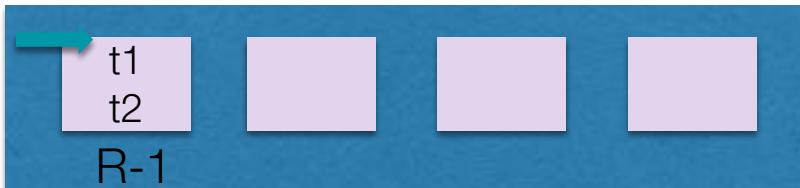
R.close()

SELECT * FROM R

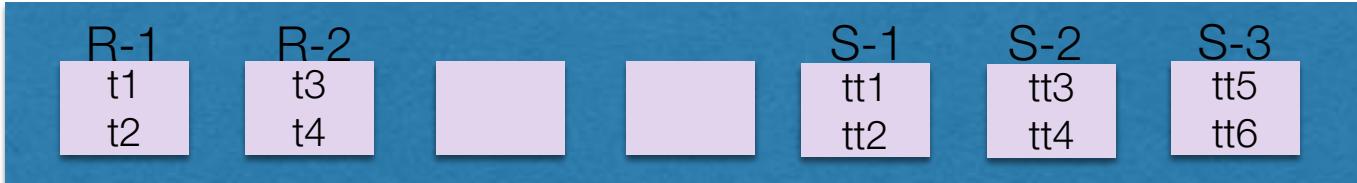
DB

Buffer

Disco



```
R.open()  
  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
  
R.close()
```

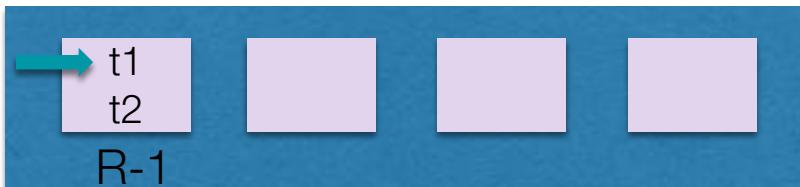


SELECT * FROM R

DB

Buffer

Disco



```
R.open()
```

```
t:= R.next()  
while t != null do  
    output t  
    t:= R.next()
```

```
R.close()
```

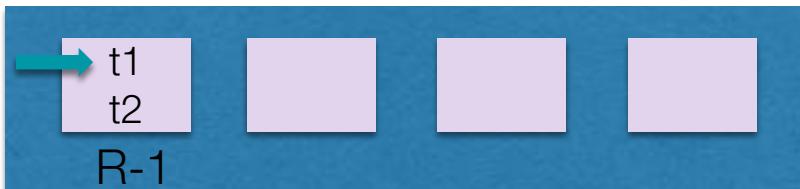


SELECT * FROM R

DB

Buffer

Disco



```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

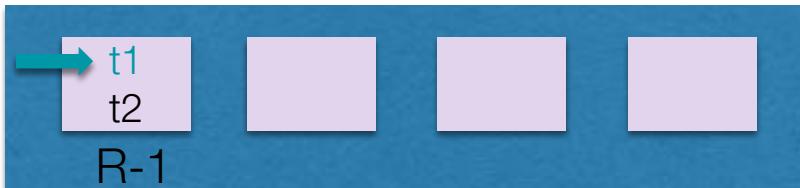


SELECT * FROM R

DB

Buffer

Disco



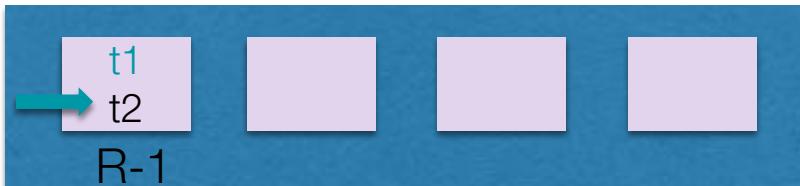
```
R.open()  
  
t:= R.next()  
while t != null do  
    output t  
    t:= R.next()  
  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco



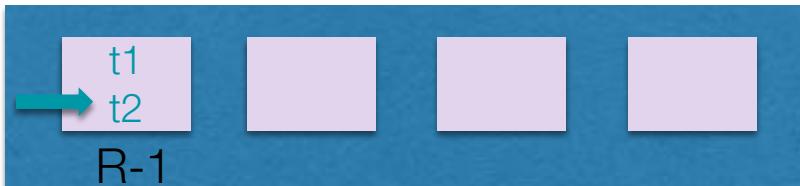
```
R.open()  
  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco



```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

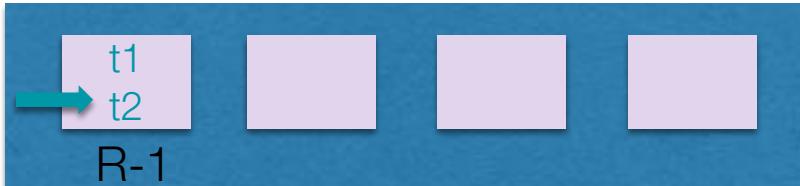


SELECT * FROM R

DB

Buffer

Disco



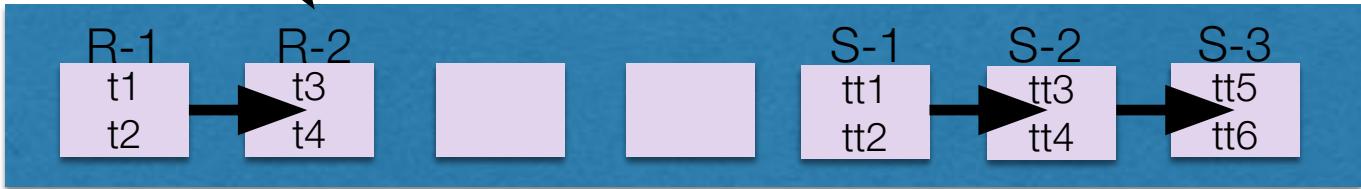
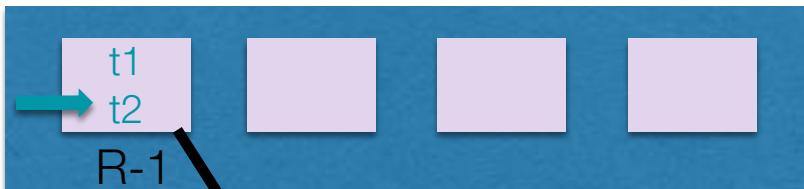
```
R.open()  
  
t:= R.next()  
while t != null do  
    output t  
    t:= R.next()  
  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco



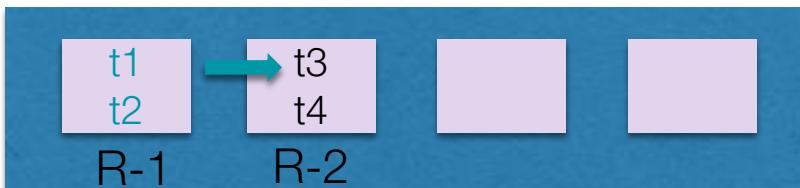
```
R.open()  
  
t:= R.next()  
while t != null do  
    output t  
    t:= R.next()  
  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco



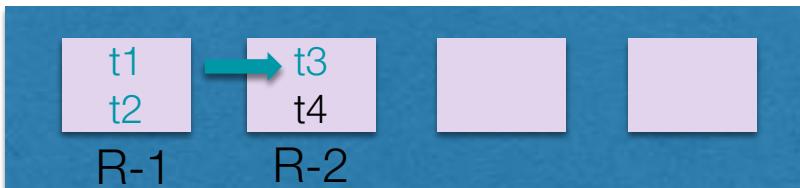
```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco



```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

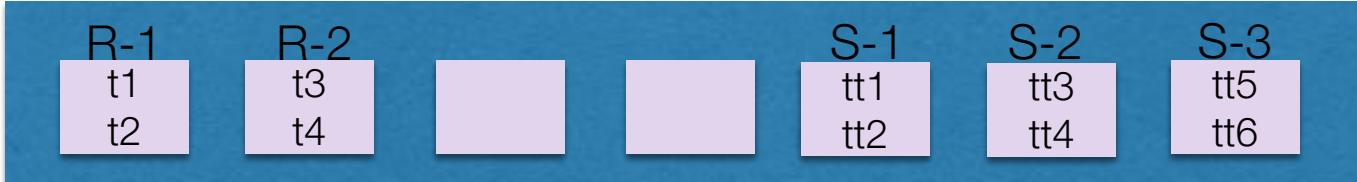
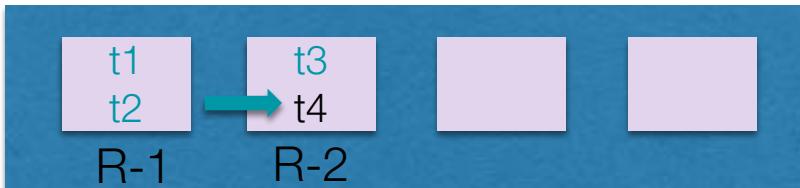


SELECT * FROM R

DB

Buffer

Disco



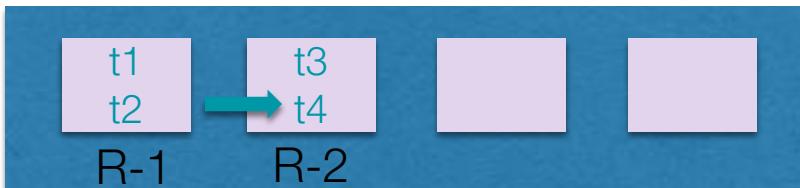
```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco



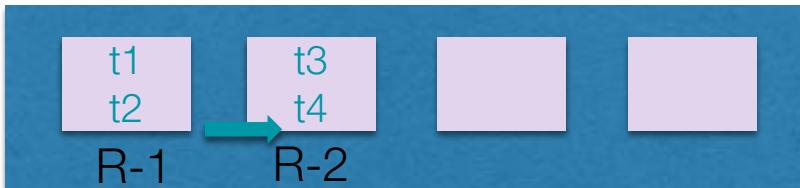
```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

SELECT * FROM R

DB

Buffer

Disco

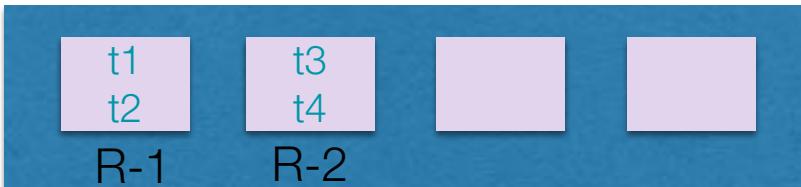


```
R.open()  
t := R.next()  
while t != null do  
    output t  
    t := R.next()  
R.close()
```

SELECT * FROM R

DB

Buffer

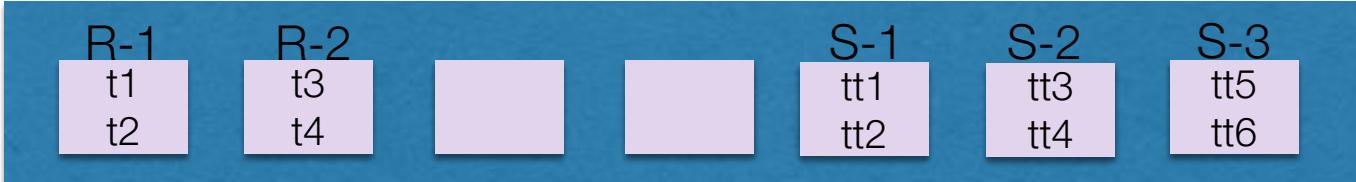


R.open()

```
t := R.next()  
while t != null do  
    output t  
    t := R.next()
```

R.close()

Disco



En realidad

Cada operador de álgebra relacional implementa interfaz de un iterador lineal:

- `open()`
- `next()`
- `close()`

Selección

El algoritmo de selección cambia dependiendo si es una consulta de igualdad (=) o de rango (<, >)

También depende si el atributo a seleccionar está indexado

Implementa interfaz de iterador lineal

Selección

Sin índice

Si queremos hacer una selección sobre una tabla **R**

```
open()  
R.open()
```

`next()` // retorna el siguiente seleccionado

```
t:= R.next()  
while t != null do  
  if t satisface condición then  
    return t  
  t:= R.next()  
return null
```

Selección

Sin índice

Si queremos hacer una selección sobre una tabla **R**

open()

R.open()

next() // retorna el siguiente seleccionado

t := R.next()

while t != null do

if t satisface condición then

return t

t := R.next()

return null

close()

R.close()

Para recorrer la selección $\text{Sel} = \sigma_{\text{cond}} (\mathbf{R})$

Sel.open()

t := Sel.next()

while t != null do

output t

t := Sel.next()

Sel.close()

Selección

Sin índice

Necesariamente tenemos que recorrer todo **R**

Selección

Con índice y consulta de igualdad

Si queremos hacer una selección sobre una tabla **R** a un atributo indexado con un índice **I**

```
open()  
I.open()  
I.search(Atributo = valor)
```

```
next()  
t:= I.next()  
while t != null do  
    return t  
return null
```

```
close()  
I.close()
```

Selección

Con índice y consulta de igualdad

Sólo tenemos que leer las páginas que satisfacen la condición
(más I/O si muchas tuplas satisfacen la condición)

Cambia un poco si el índice es Clustered o Unclustered (¿Por qué?)

Si el atributo es llave primaria entonces la operación prácticamente tiene I/O ~ 1

Selección

Con índice y consulta de rango

¿Cómo podemos hacer este tipo de consultas de forma eficiente?

Hint Existe un índice especial para hacer esto

Proyección

Algoritmo muy sencillo

```
open()  
  R.open()
```

```
next()  
  t:= R.next()  
  while t != null do  
    return project(t, atributos)  
  return null
```

```
close()  
  R.close()
```

Proyección

Necesariamente tenemos que recorrer todo **R**

Joins

Operación muy costosa

Supondremos solamente restricciones de igualdad (por ejemplo, $R.a = S.a$)

Nested Loop Join

Queremos hacer un join entre **R** y **S**, cuando se satisface un predicado **p**

```
open()  
  R.open()  
  S.open()  
  r:= R.next()
```

```
close()  
  R.close()  
  S.close()
```

Nested Loop Join

Queremos hacer un join entre **R** y **S**, cuando se satisface un predicado **p**

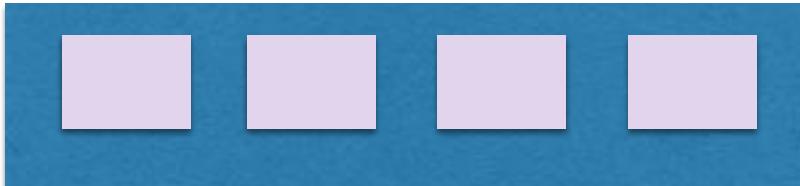
```
next()
  while r != null do
    s:= S.next()
    if s == null then
      S.close()
      r:= R.next()
      S.open()
    else if (r, s) satisfacen p then
      return (r, s)
  return null
```

Nested Loop Join

DB

Buffer

Disco



R-1	R-2	R-3	R-4	S-1	S-2	S-3
t1 t2	t3 t4	t5 t6	t7 t8	tt1 tt2	tt3 tt4	tt5 tt6

$$\text{Join} = R \bowtie_p S$$

Join.open()

```
t := Join.next()  
while t != null do  
    output t  
    t := Join.next()
```

Join.close()

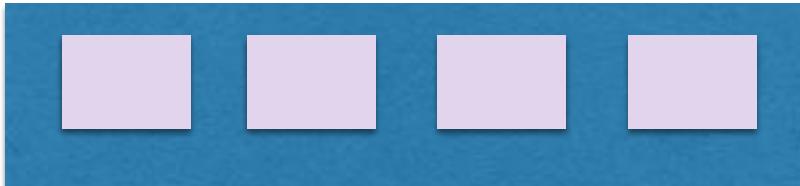
Nested Loop Join

DB

$R \bowtie_p S \dots \text{Join.open}()$

`R.open()
S.open()
r := R.next()`

Buffer



Disco

	R-1	R-2	R-3	R-4	S-1	S-2	S-3
	t1 t2	t3 t4	t5 t6	t7 t8	tt1 tt2	tt3 tt4	tt5 tt6

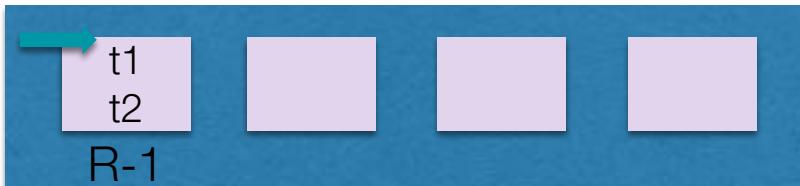
Nested Loop Join

DB

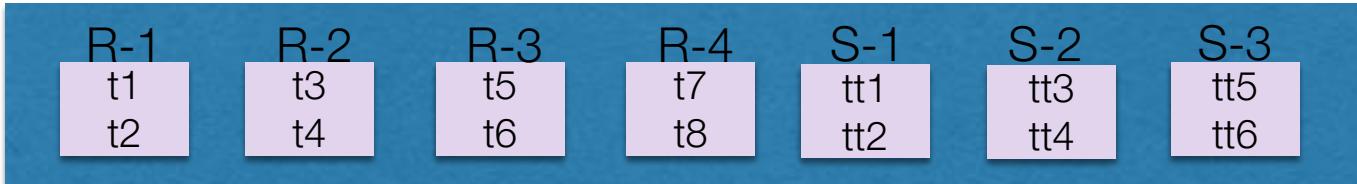
$R \bowtie_p S \dots \text{Join.open}()$

`R.open()
S.open()
r := R.next()`

Buffer



Disco



Nested Loop Join

DB

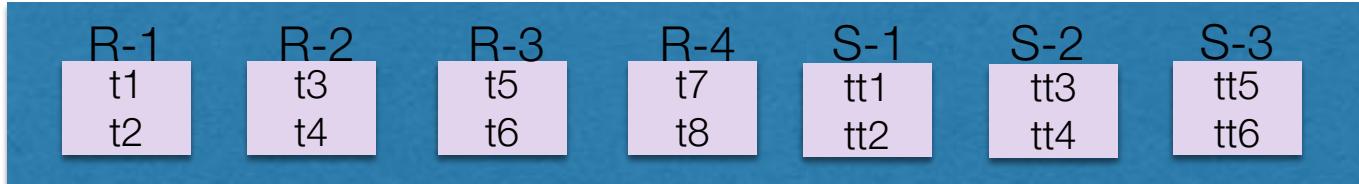
$R \bowtie_p S \dots \text{Join.open}()$

`R.open()
S.open()
r := R.next()`

Buffer



Disco



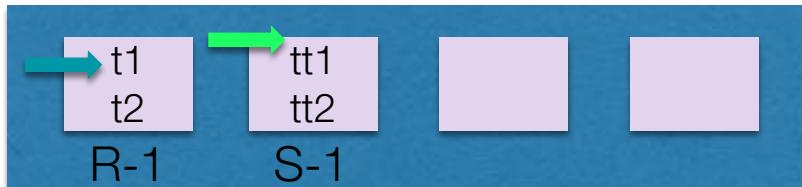
Nested Loop Join

DB

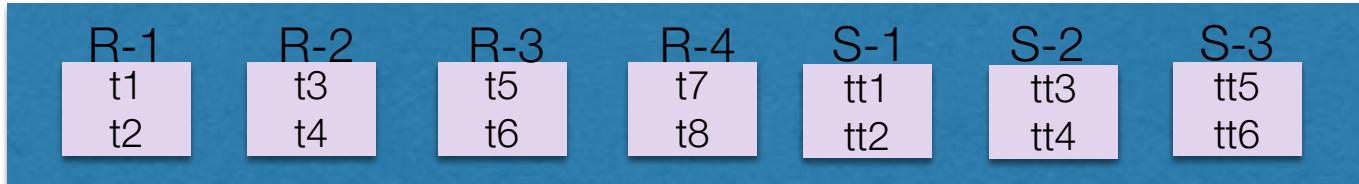
$R \bowtie_p S \dots \text{Join.open}()$

`R.open()
S.open()
r := R.next()`

Buffer



Disco

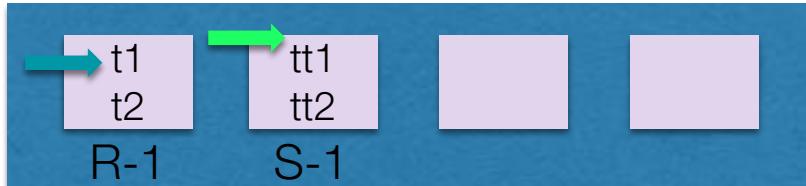


Nested Loop Join

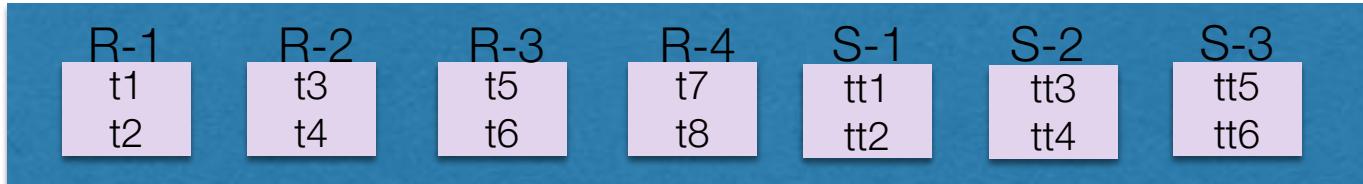
DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

Buffer



Disco



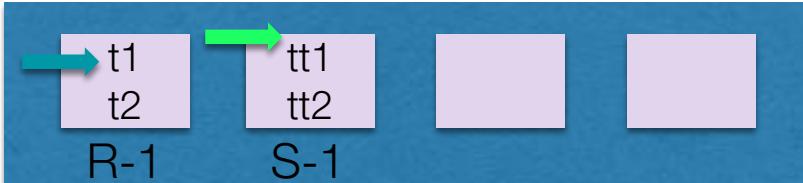
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



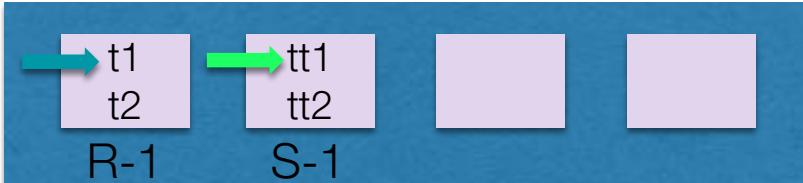
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



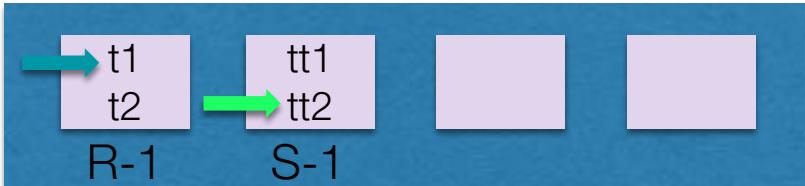
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



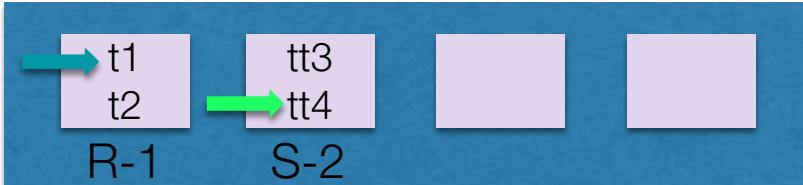
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



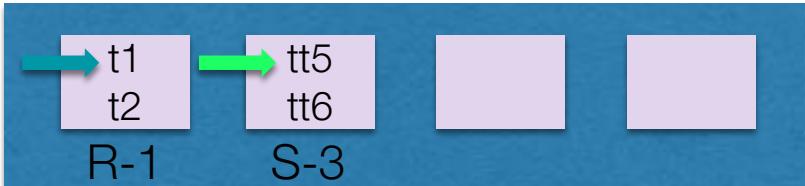
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



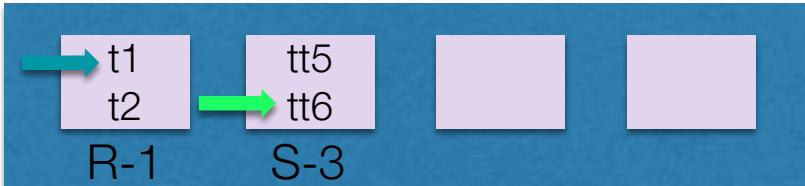
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco

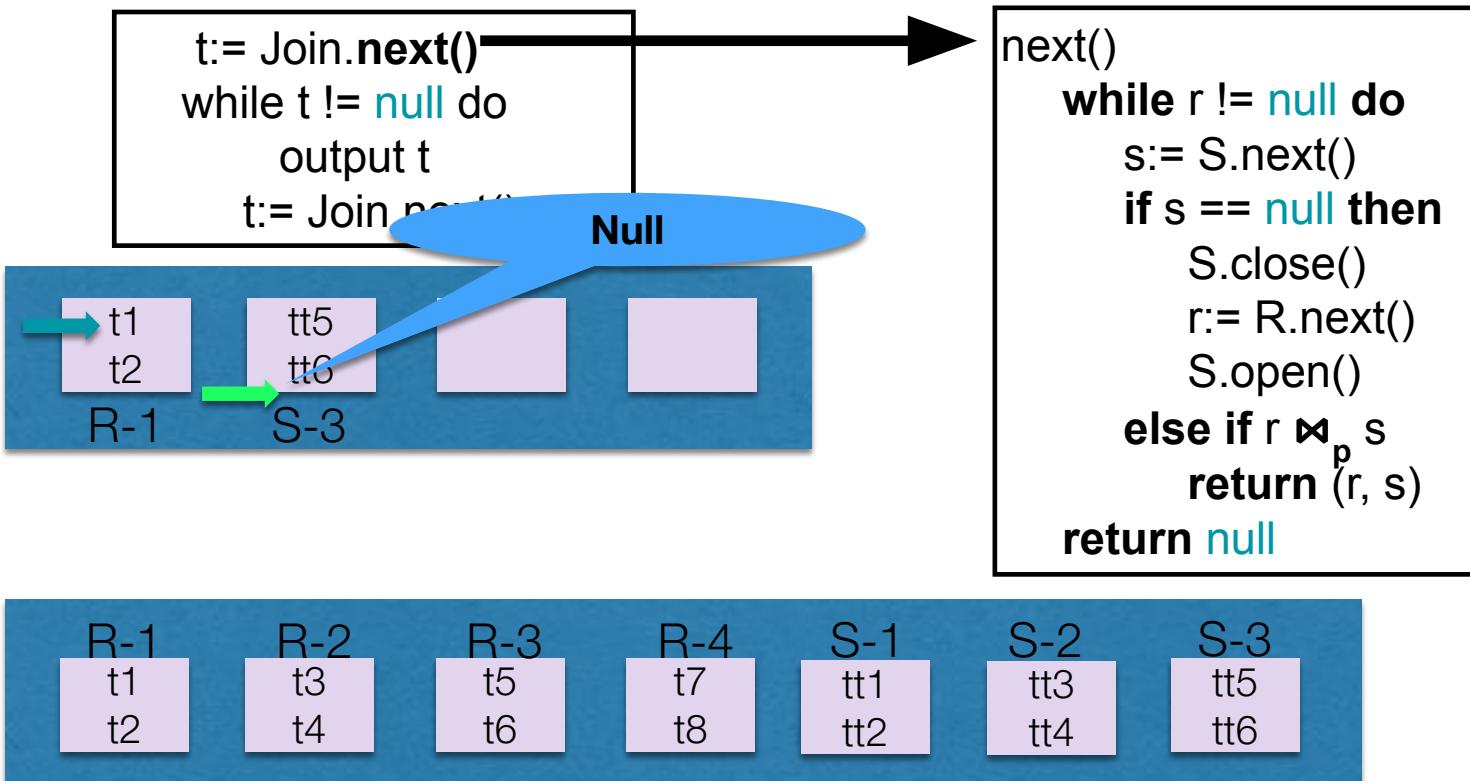


Nested Loop Join

DB

Buffer

Disco



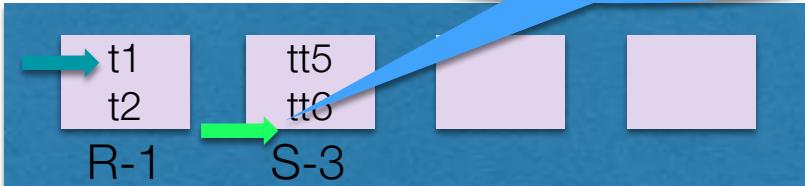
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco

	R-1	R-2	R-3	R-4	S-1	S-2	S-3
	t1 t2	t3 t4	t5 t6	t7 t8	tt1 tt2	tt3 tt4	tt5 tt6

Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



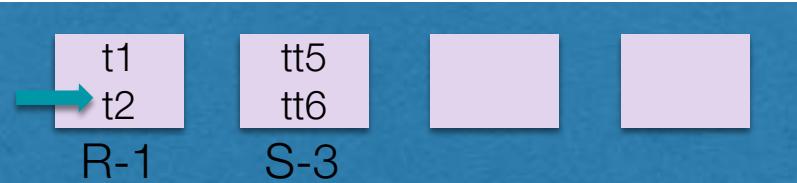
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



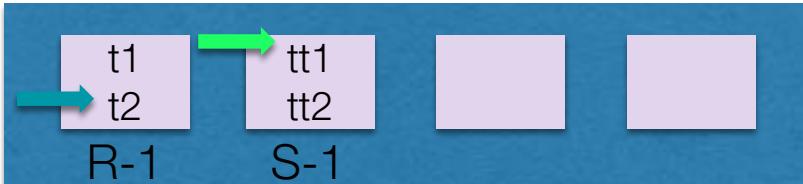
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco



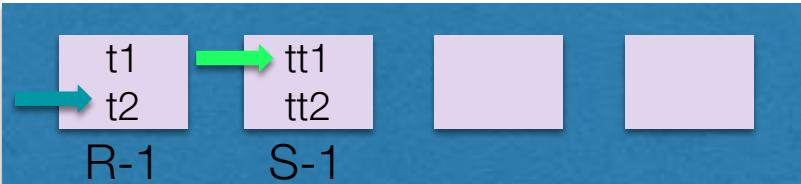
Nested Loop Join

DB

```
t:= Join.next()  
while t != null do  
    output t  
    t:= Join.next()
```

```
next()  
while r != null do  
    s:= S.next()  
    if s == null then  
        S.close()  
        r:= R.next()  
        S.open()  
    else if r  $\bowtie_p$  s  
        return (r, s)  
    return null
```

Buffer



Disco

R-1	R-2	R-3	R-4	S-1	S-2	S-3
t1	t3	t5	t7	tt1	tt3	tt5
t2	t4	t6	t8	tt2	tt4	tt6

Nested Loop Join

Es una implementación directa basada en un loop

Para cada tupla de **R** debemos leer **S** entera, aparte de leer **R** entera una vez

Costo en I/O es:

$$\text{Costo}(\mathbf{R}) + \text{Tuplas}(\mathbf{R}) \cdot \text{Costo}(\mathbf{S})$$

Nested Loop Join

Si **R** y **S** son tablas de 16 MB, cada página es de 8 KB y las tuplas son de 300 bytes

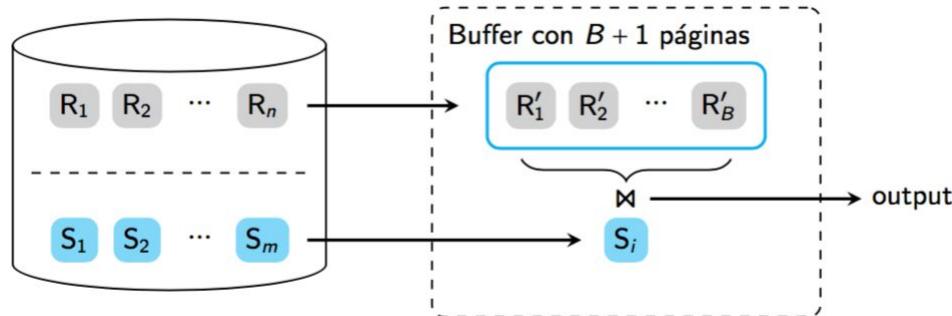
Cada relación tiene 2048 páginas y 55.000 tuplas aproximadamente

Costo de un I/O es 0.1 ms, entonces el join tarda:

3.1 horas

Block Nested Loop Join

Aprovechamos mejor el buffer



Block Nested Loop Join

Queremos hacer un join entre **R** y **S**, cuando se satisface un predicado **p**

```
open()  
  R.open()  
  fillBuffer()
```

```
close()  
  R.close()  
  S.close()
```

```
fillBuffer()  
  Buff = empty  
  r := R.next()  
  while r != null do  
    Buff = Buff union r  
    if Buff.isFull() then  
      break  
    r := R.next()  
  S.open()  
  S.next()
```

Block Nested Loop Join

Queremos hacer un join entre **R** y **S**, cuando se satisface un predicado **p**

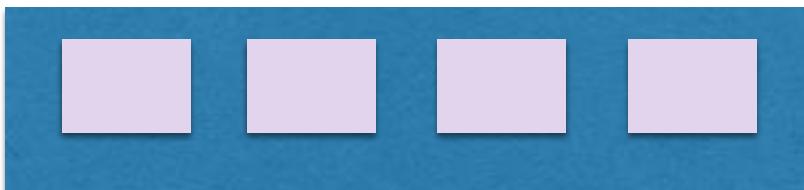
```
next()
  while Buff != empty do
    while s != null do
      r:= Buffer.next()
      if r == null then
        Buffer.reset()
        s:= S.next()
      else if (r,s) satisfacen p then
        return (r,s)
      fillBuffer()
  return null
```

Block Nested Loop Join

DB

 $R \bowtie_p S$

Buffer



Disco

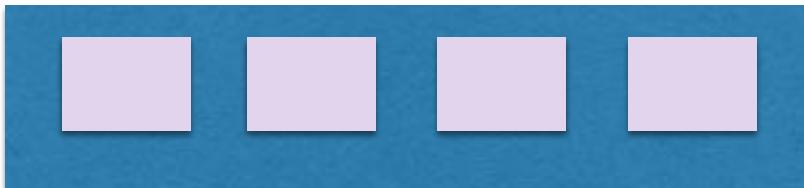


Block Nested Loop Join

DB

$R \bowtie_p S.\text{open}()$

Buffer



Disco

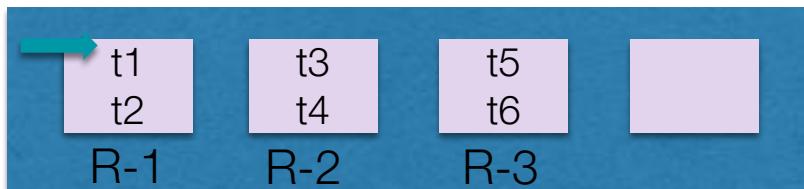


Block Nested Loop Join

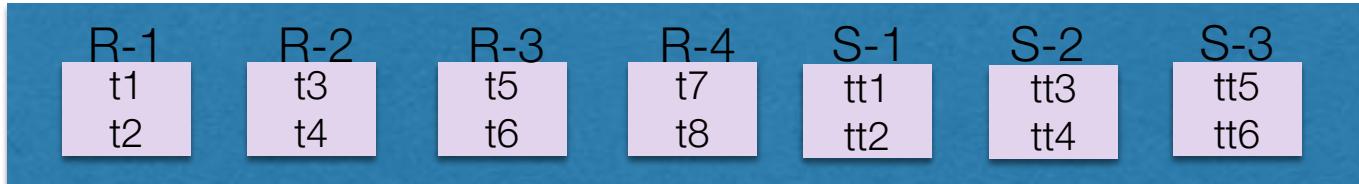
DB

$R \bowtie_p S.\text{open}()$

Buffer



Disco

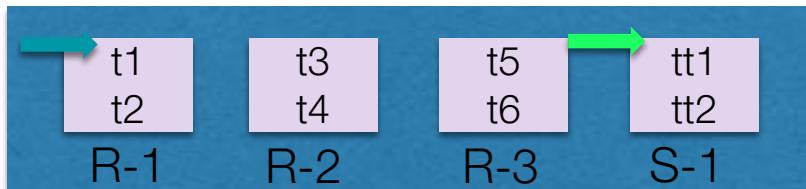


Block Nested Loop Join

DB

$R \bowtie_p S.\text{open}()$

Buffer



Disco

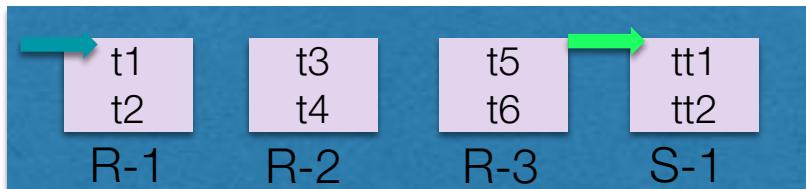


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

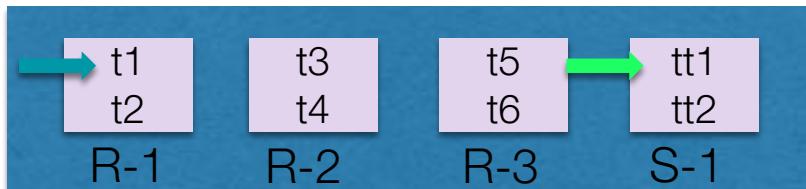


Block Nested Loop Join

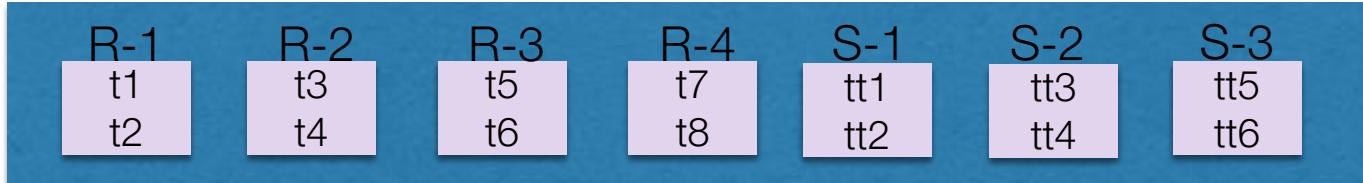
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

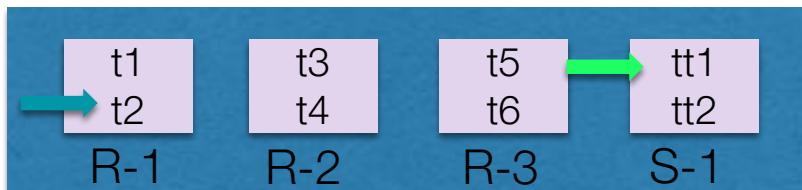


Block Nested Loop Join

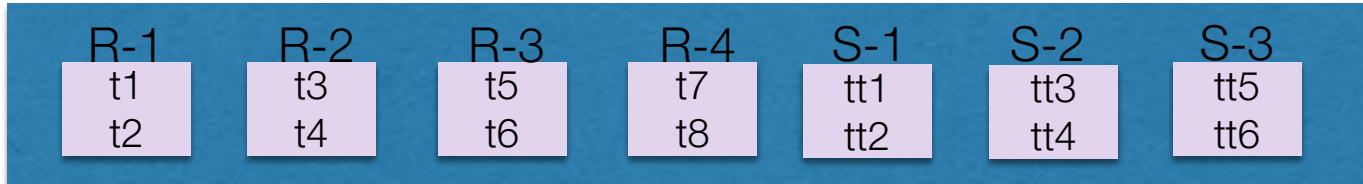
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

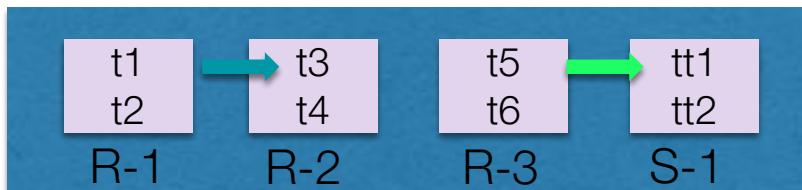


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

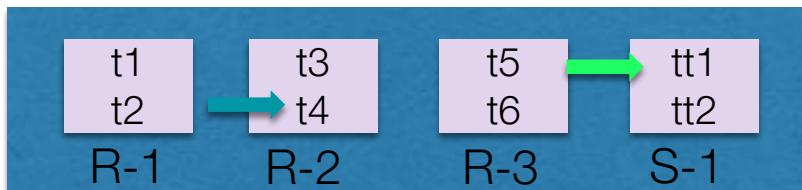


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

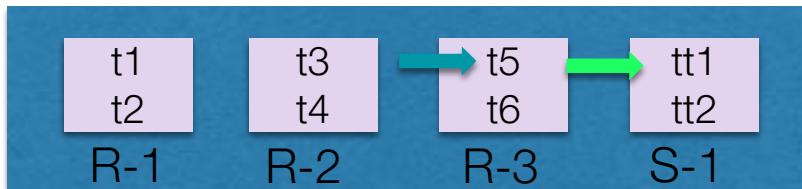


Block Nested Loop Join

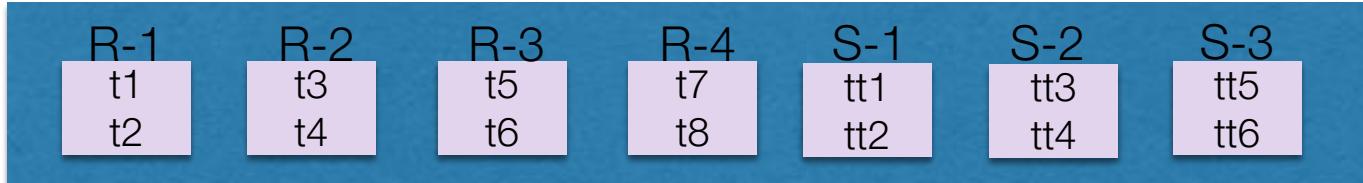
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

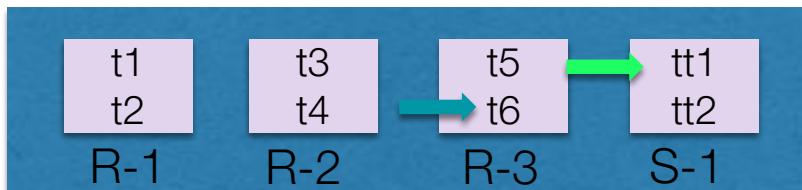


Block Nested Loop Join

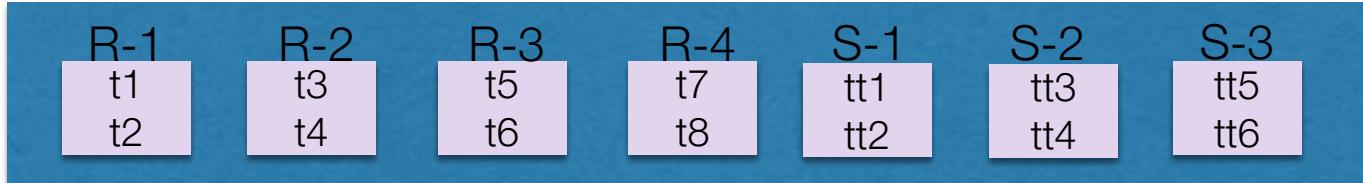
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco



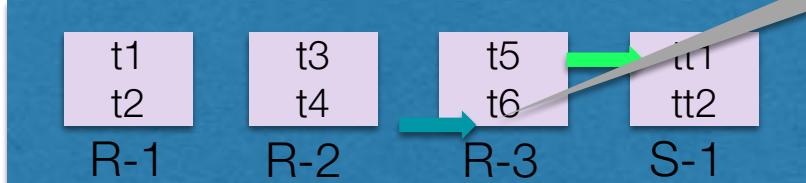
Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

End of Buffer de R

Buffer



Disco

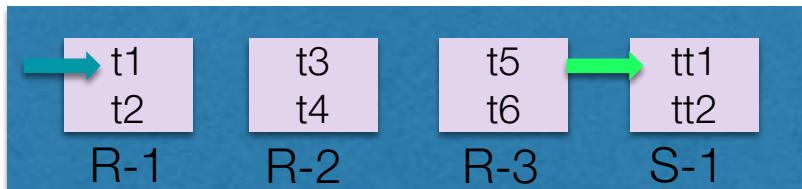


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

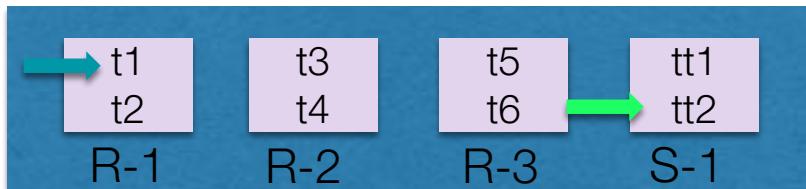


Block Nested Loop Join

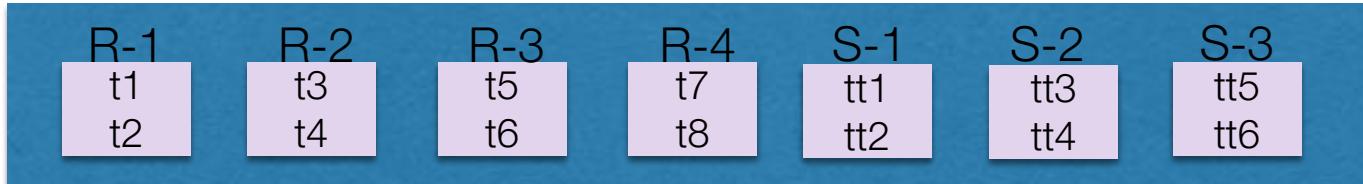
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

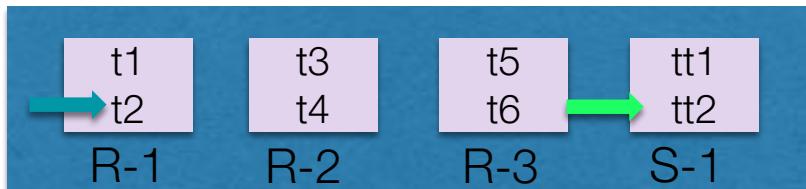


Block Nested Loop Join

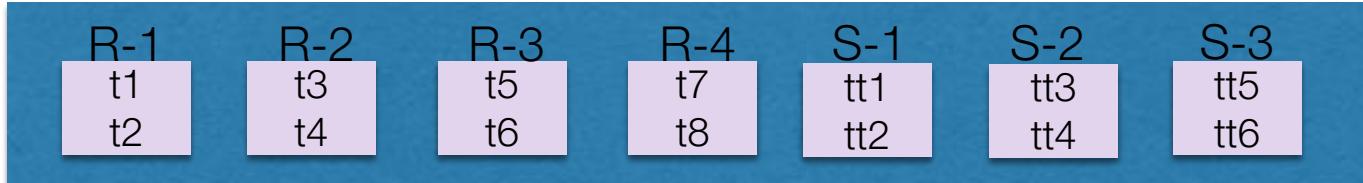
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

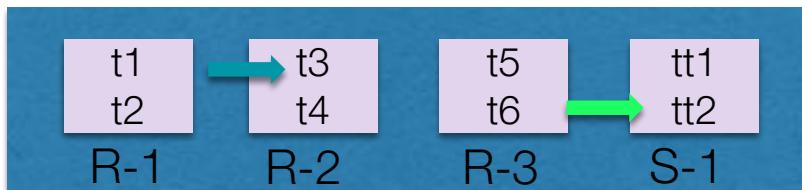


Block Nested Loop Join

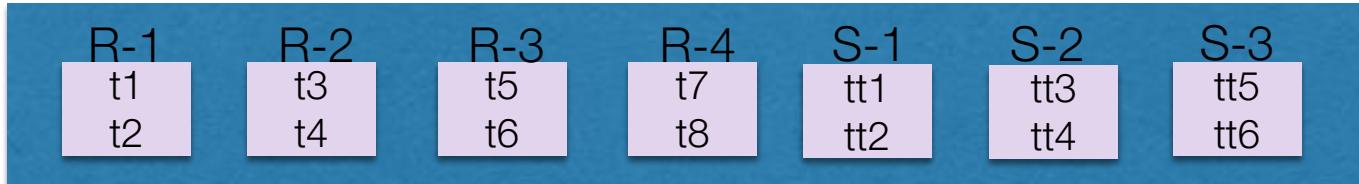
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

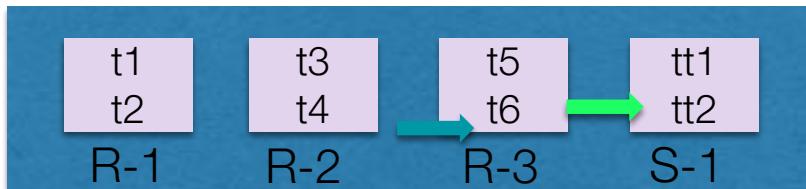


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

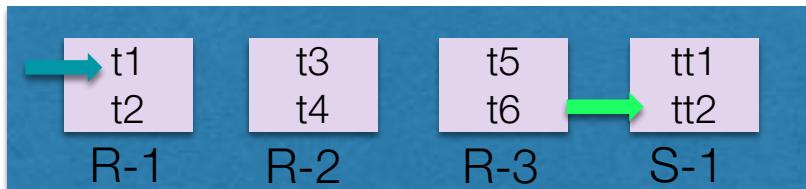


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

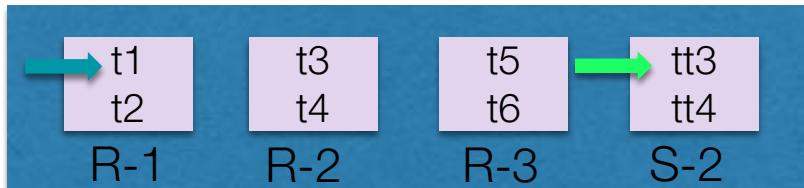


Block Nested Loop Join

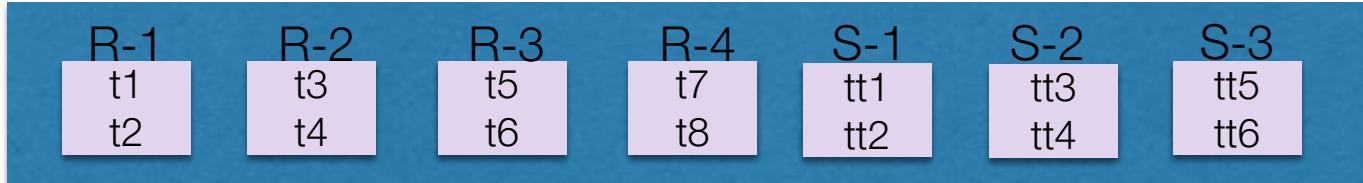
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

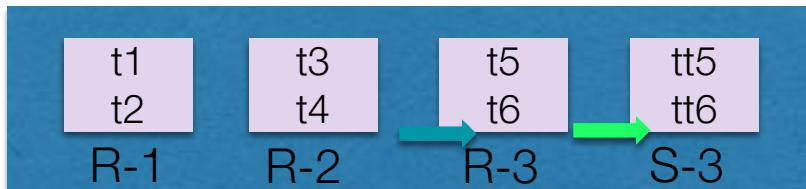


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

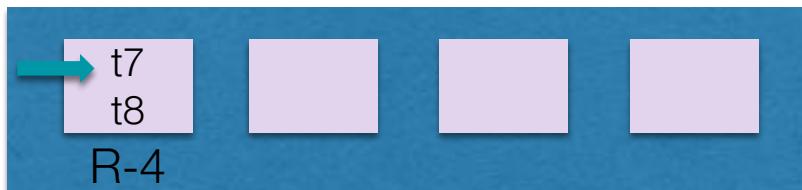


Block Nested Loop Join

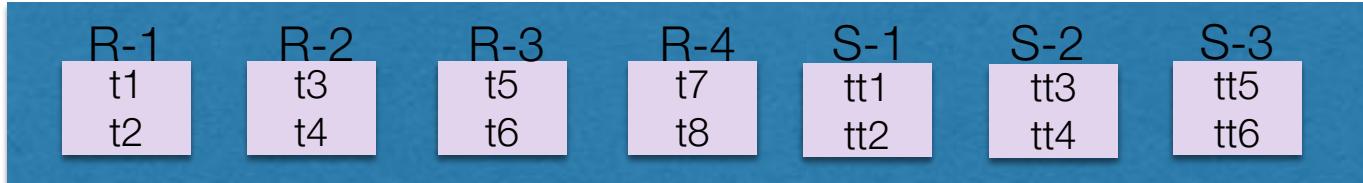
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

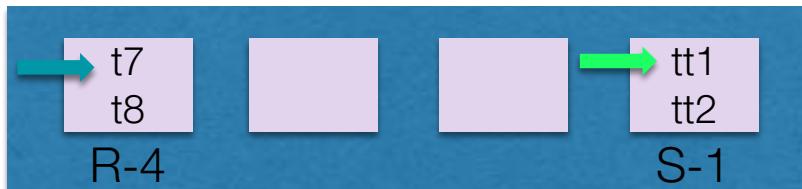


Block Nested Loop Join

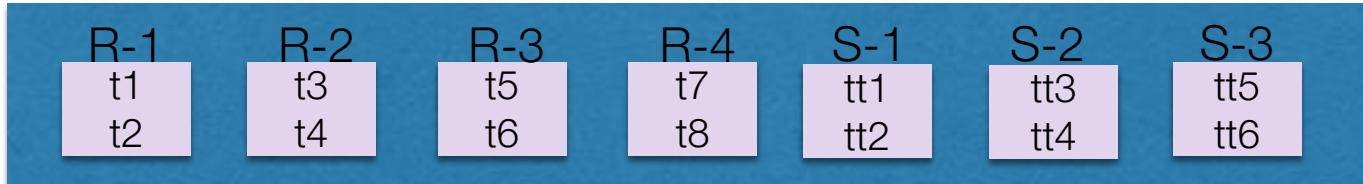
DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco

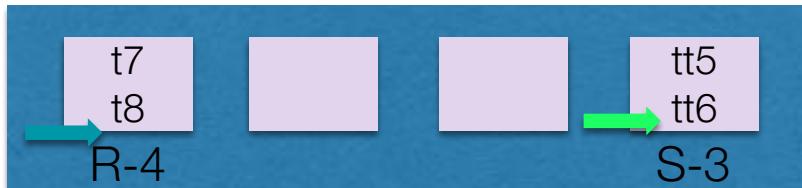


Block Nested Loop Join

DB

```
while (R  $\bowtie_p$  S.next())
```

Buffer



Disco



Block Nested Loop Join

Ahora cargamos muchas páginas de **R** a buffer

Por cada vez que llenamos el buffer recorremos a **S** entera una vez

Costo en I/O es:

$$\text{Costo}(\mathbf{R}) + (\text{Páginas}(\mathbf{R})/\text{Buffer}) \cdot \text{Costo}(\mathbf{S})$$

Block Nested Loop Join

Si **R** y **S** son tablas de 16 MB, cada página es de 8 KB con un **buffer** de 1 MB

Cada relación tiene 2048 páginas y en **buffer** caben 128 páginas

Costo de un I/O es 0.1 ms, entonces el join tarda:

3.4 segundos

Block Nested Loop Join

Sin embargo existen algoritmos muchos más eficientes

Estos algoritmos se basan en Hashing o en Sorting

Además hacen usos de índices, como por ejemplo el B+ Tree

Sorting

Los algoritmos de sorting son conocidos en programación

¿Por qué estudiarlos otra vez?

Sorting

Necesitamos ordenar tuplas que exceden por mucho el tamaño de la memoria RAM

External Merge Sort

En los DBMS, se utiliza el algoritmo External Merge Sort

Hablaremos de **Run** como una secuencia de páginas que contiene una conjunto ordenado de tuplas

Algoritmo funciona por **fases**

External Merge Sort

Fase 0: creamos los runs iniciales

Fase i:

- Traemos los runs a memoria
- Hacemos el merge de cada par de runs
- Almacenamos el nuevo run a disco (i.e. materializamos resultados intermedios)

External Merge Sort

